

# **Count Response**

STAT 526  
Professor Olga Vitek

March 31, 2011


# **Basic Poisson Regression (Faraway Ch. 3)**

# Poisson Distribution


- $Y$ : Counts of events that occur randomly over fixed interval of time or space
  - No upper limit for the number of events
  - Approximates  $Binomial(n, \pi)$  distribution when  $n$  is large and  $\pi$  is small. Then  $\lambda = n\pi$ .
  - Approaches Normality as  $\lambda$  increases
- $P\{Y = y\} = \frac{e^{-\lambda} \lambda^y}{y!}, \quad y = 1, \dots$
- $E\{Y\} = Var\{Y\} = \lambda$
- If  $Y_i \stackrel{ind}{\sim} Poisson(\lambda_i), \quad i = 1, \dots,$   
then  $\sum_i Y_i \sim Poisson(\sum_i \lambda_i)$

# Poisson( $\lambda$ ) for various $\lambda$


$\lambda = 0.5$




$\lambda = 1$




$\lambda = 3$




$\lambda = 6$



$\lambda = 9$



$\lambda = 12$



# Poisson Regression

- Model

$$Y_i \stackrel{ind}{\sim} \text{Poisson}(\lambda_i), \text{ where} \\ \log \lambda_i = \mathbf{X}_i' \beta \quad (= \eta, \text{ link})$$

–  $\beta_j$  = difference in  $\log E\{Y\}$  following a unit change in predictor  $X_j$ , while the others are constant

- Log-likelihood

$$\begin{aligned} l(\beta) &= \log \prod_{i=1}^n \left[ \frac{\lambda_i^{y_i} e^{-\lambda_i}}{y_i!} \right] = \sum_{i=1}^n (y_i \log \lambda_i - \lambda_i - \log y_i!) \\ &= \sum_{i=1}^n (y_i \mathbf{X}_i' \beta - \exp \mathbf{X}_i' \beta) + \text{constant} \end{aligned}$$

- Pearson  $X^2$

$$X^2 = \sum_{i=1}^n \frac{(y_i - \hat{\lambda}_i)^2}{\hat{\lambda}_i}$$

- Deviance

$$D = 2 \sum_{i=1}^n \left( y_i \log \frac{y_i}{\hat{\lambda}_i} - (y_i - \hat{\lambda}_i) \right)$$

– Different numeric values of deviance for grouped and ungrouped data

# Inference About Individual $\beta_j$ : Wald Test

- Same as for Binomial response
- Test  $H_0 : \beta_j = 0$  versus  $H_a : \beta_j \neq 0$ .
- Test statistic  $z^* = \frac{b_j - 0}{s\{b_j\}}$
- **Approximate** variance  $s^2\{\mathbf{b}\}$

$$s^2\{\mathbf{b}\} = \left( \left[ - \frac{\partial^2 \log_e L(\beta)}{\partial \beta_j \partial \beta_{j'}} \right]_{\beta=\mathbf{b}} \right)^{-1}$$

- **Approximate** distribution of  $z$ 
  - $z^* \sim \mathcal{N}(0, 1)$ . Alternatively,  $(z^*)^2 \sim \chi_1^2$
  - reject  $H_0$  if  $|z^*| > z^{1-\alpha/2}$
  - CI for  $\beta_j$ :  $b_j \pm z^{1-\alpha/2} s\{b_j\}$

# Example: Galapagos Plants

## (Faraway Ch. 3)

- Study the number of plant species on 30 Galapagos islands as function of geographic variables

```
library(faraway)
data(gala)
?gala
```

```
# faraway does not use the 2nd predictor,
# so we'll remove it too.
gala <- gala[, -2]
head(gala)
```

	Species	Area	Elevation	Nearest	Scruz	Adjacent
Baltra	58	25.09	346	0.6	0.6	1.84
Bartolome	31	1.24	109	0.6	26.3	572.33
Caldwell	3	0.21	114	2.8	58.7	0.78
Champion	25	0.10	46	1.9	47.4	0.18
Coamano	2	0.05	77	1.9	1.9	903.82
Daphne.Major	18	0.34	119	8.0	8.0	1.84

# Model with all the predictors

```
> fit1 <- glm(Species~., family=poisson, data=gala)
> summary(fit1)
```

Call:

```
glm(formula = Species ~ ., family = poisson, data = gala)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-8.2752	-4.4966	-0.9443	1.9168	10.1849

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	3.155e+00	5.175e-02	60.963	< 2e-16 ***
Area	-5.799e-04	2.627e-05	-22.074	< 2e-16 ***
Elevation	3.541e-03	8.741e-05	40.507	< 2e-16 ***
Nearest	8.826e-03	1.821e-03	4.846	1.26e-06 ***
Scruz	-5.709e-03	6.256e-04	-9.126	< 2e-16 ***
Adjacent	-6.630e-04	2.933e-05	-22.608	< 2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 3510.73 on 29 degrees of freedom  
Residual deviance: 716.85 on 24 degrees of freedom  
AIC: 889.68



# **Model Diagnostics (Faraway Sec. 6.4)**

# Goodness of Fit Tests

- Residual Deviance

- Test  $H_0$  : our model vs  $H_a$ : saturated model
- Distribution poorly approximated by  $\chi^2$

```
> pchisq(fit1$deviance,  
         fit1$df.residual, lower.tail=FALSE)  
[1] 7.073157e-136 # reject H0
```

- Pearson  $\chi^2$

- Test  $H_0$  : our model vs  $H_a$ : saturated model
- Distribution better approximated by  $\chi^2$

```
> pchisq( sum(residuals(fit1, type="pearson")^2 ),  
         fit1$df.residual, lower.tail=FALSE)  
[1] 2.18719e-145 # reject H0
```

- Better approximation of  $\chi^2$  when Poisson approaches Normal
- Can look more into model diagnostics to find possible reasons for poor fit

# Model Diagnostics: Residual Plots

- Response residuals  $r_i = y_i - \hat{\lambda}_i$ 
  - Of limited use: variance is function of the expected value
- Pearson residuals  $r_{i,P} = \frac{y_i - \hat{\lambda}_i}{\sqrt{\hat{\lambda}_i}}$ 
  - $\sum_i r_{i,P}^2 = X^2$
  - Account for differences in variance
- Deviance residuals
$$r_{i,D} = \text{sign}(y_i - \hat{\lambda}_i) \left[ 2(y_i \log \frac{y_i}{\hat{\lambda}_i} - (y_i - \hat{\lambda}_i)) \right]^{1/2}$$
  - Defined by analogy with Pearson residuals
  - Define  $d_i = 2(y_i \log \frac{y_i}{\hat{\lambda}_i} - (y_i - \hat{\lambda}_i))$ , such that
$$\sum_i r_{i,D}^2 = \sum_i d_i = \text{Deviance}$$
  - Account for differences in variance

# Residual Plots (Next Slide)

(a) As in linear regression, can plot response residuals  $r_i$  vs  $E\{\widehat{Y}_i\} = \widehat{\lambda}_i$

```
plot(residuals(fit1, type="response") ~  
     predict(fit1, type="response"),  
      xlab=expression(hat(lambda)),  
      ylab="Response residuals", main="(a)")  
abline(h=0)
```

- Of limited use:  
Variance is function of the expected value;  
Most predicted expected values are small

(b) Plot response residuals  $r_i$  vs link  $\widehat{\eta} = X\widehat{\beta}$

```
plot(residuals(fit1, type="response") ~  
     predict(fit1, type="link"),  
      xlab=expression(paste(hat(eta), " = X", hat(beta))),  
      ylab="Response residuals", main="(b)")  
abline(h=0)
```

- The points on x axis are better separated  
Variance still function of the expected value

# Residual Plots (Next Slide)

(c) Plot deviance residuals  $r_{i,D}$  vs link  $\hat{\eta} = X\hat{\beta}$

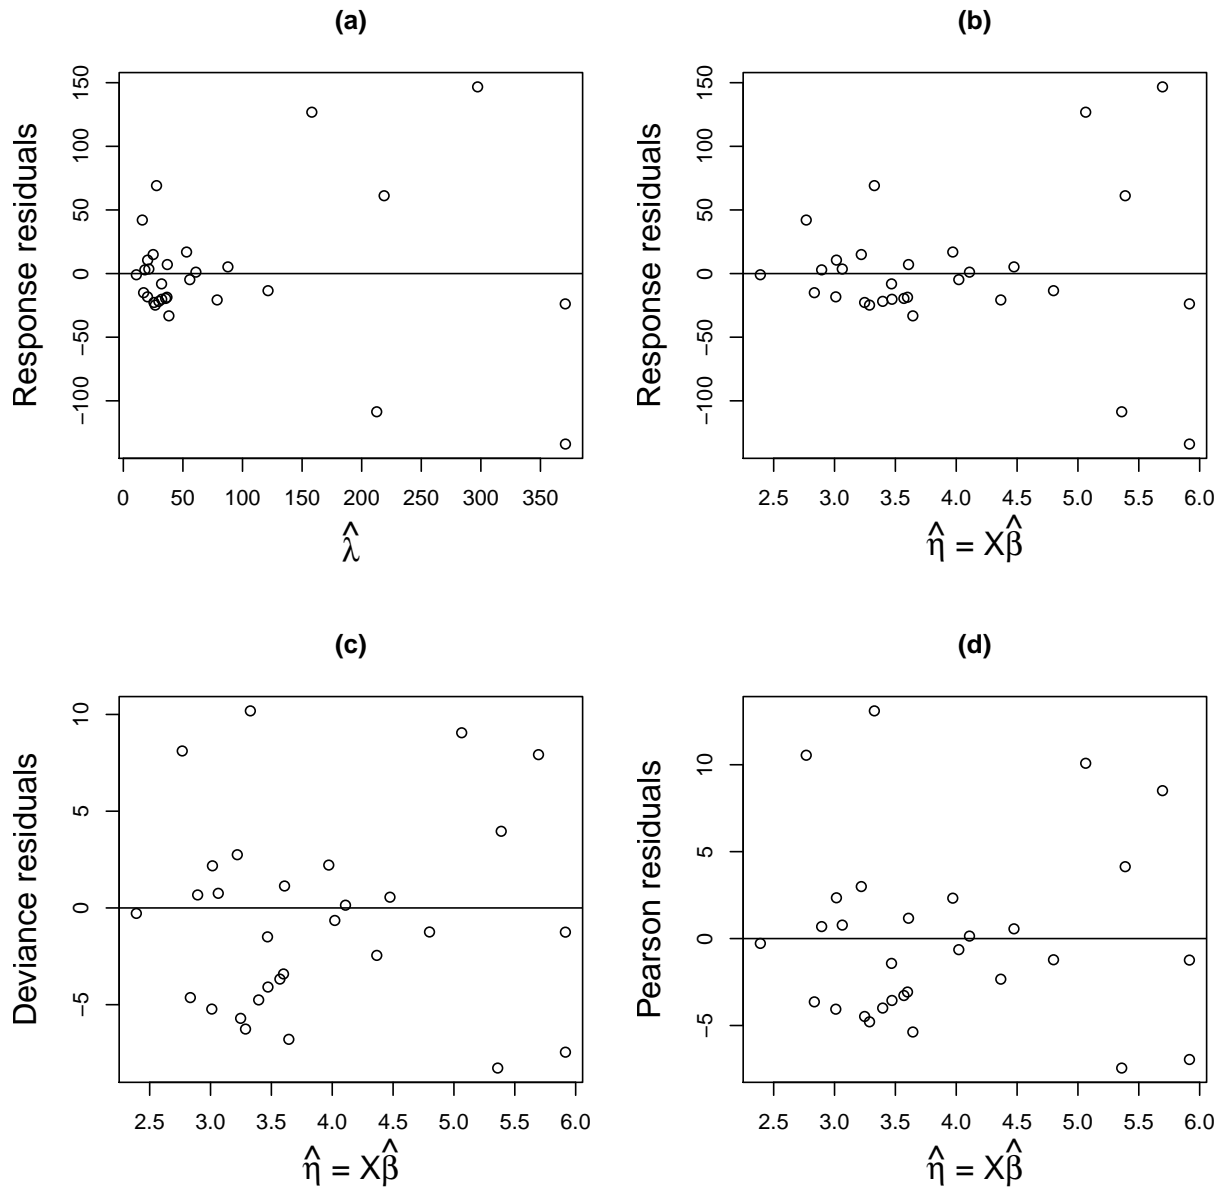
```
plot(residuals(fit1) ~
      predict(fit1, type="link"),
      xlab=expression(paste(hat(eta), " = X", hat(beta))),
      ylab="Deviance residuals", main="(c)")
abline(h=0)
```

(d) Plot Pearson residuals  $r_{i,P}$  vs link  $\hat{\eta} = X\hat{\beta}$

```
plot(residuals(fit1, type="pearson") ~
      predict(fit1, type="link"),
      xlab=expression(paste(hat(eta), " = X", hat(beta))),
      ylab="Pearson residuals", main="(d)")
abline(h=0)
```

- Interpretation of (c) and (d):
  - Y axis accounts for differences in variance
  - X axis better separates the predicted values
  - Systematic nonlinear relationship → lack of fit
    - \* No apparent problems for this dataset
  - Large/nonconstant variance → violations of the assumed Poisson distribution
    - \* Potential overdispersion

# Residual Plots



- No obvious problematic trends in (c) or (d)

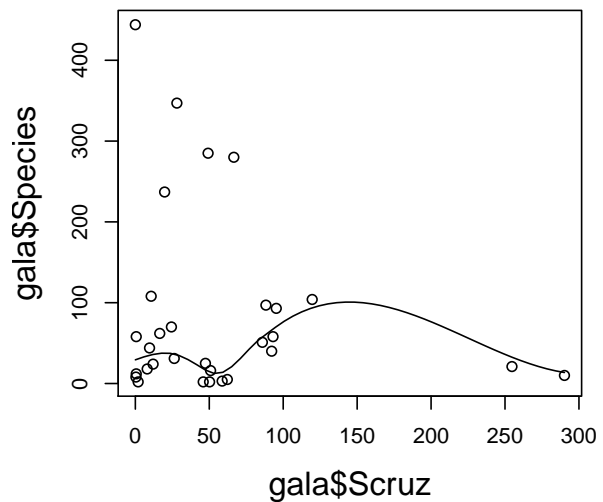
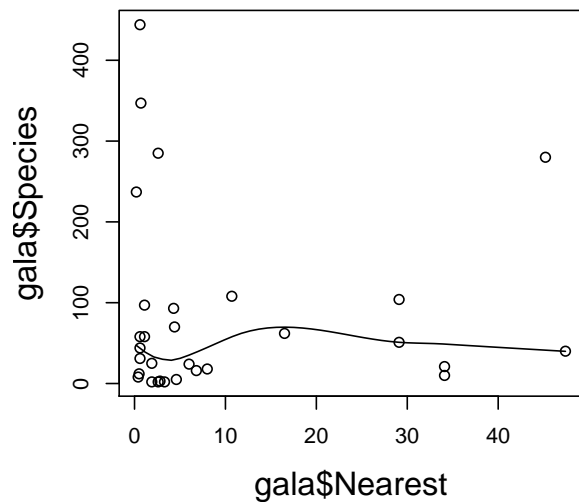
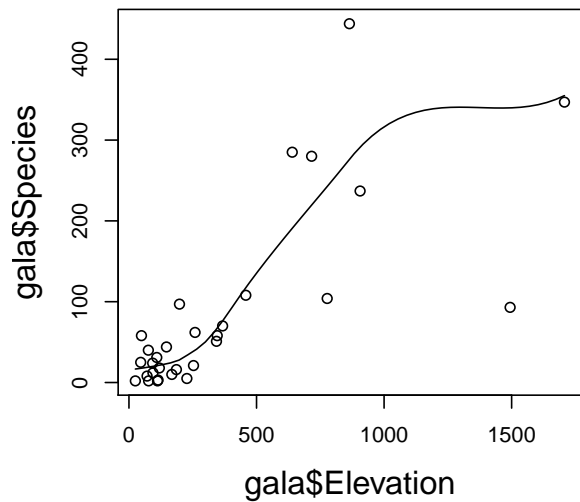
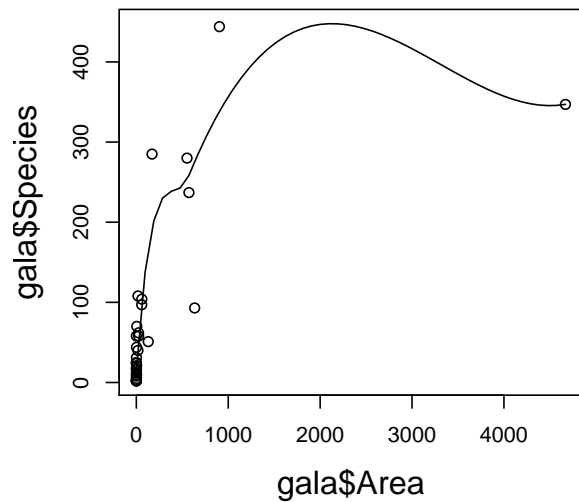
# Model Diagnostics: Functional Form of Predictors

- As in linear regression, can plot  $Y$  against each predictor  $X_j$

```
scatter.smooth(gala$Area, gala$Species)
scatter.smooth(gala$Elevation, gala$Species)
scatter.smooth(gala$Nearest, gala$Species)
scatter.smooth(gala$Scruz, gala$Species)
```

- The plots are of limited use, and may not show clear trends:
  - Lower response counts are more frequent
  - $Var\{Y\} = E\{Y\} \rightarrow$  large  $Y$  are unreliable
- Transformations can be helpful
  - Look for transformations for predictors even if there is no obvious problems in residual plots

# Plotting $Y$ vs $X_j$



- Most predictors indicate some non-linearity



## Consider $\log(\text{Area})$

(a) *Area* vs *Species*

```
plot(gala$Area, gala$Species, main="(a)")
```

(b) Use  $\log(\text{Area})$

```
scatter.smooth(log(gala$Area), gala$Species,  
               main="(b)")
```

(c) Account for the  $\log$  link:

$\log(\text{Area})$  vs *Species*

```
scatter.smooth(log(gala$Area), log(gala$Species),  
               main="(c)")
```

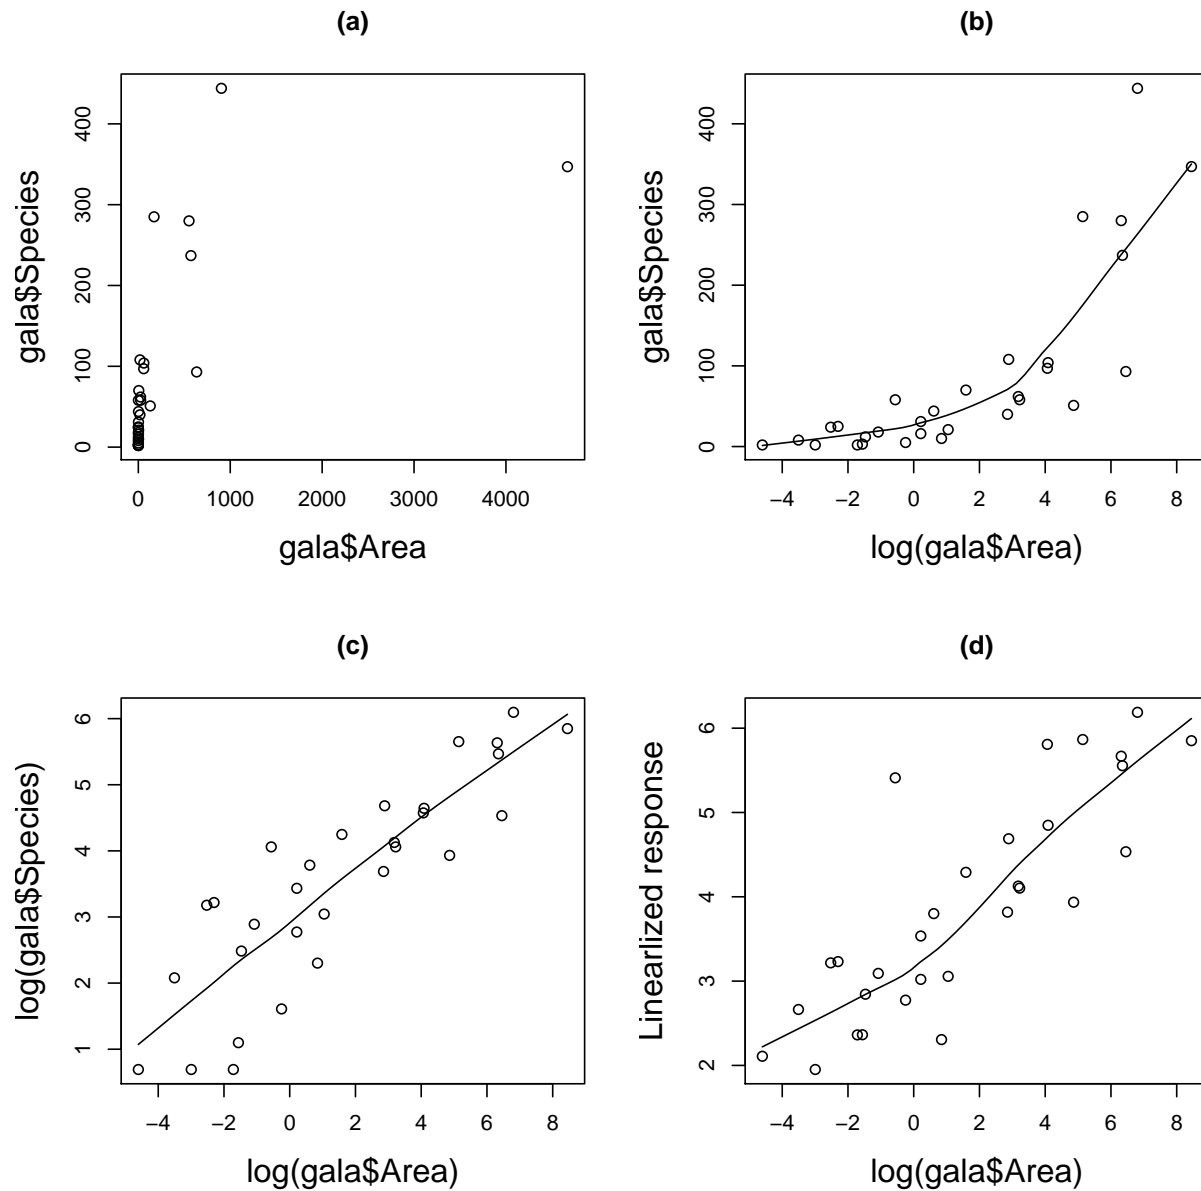
(d) Faraway suggests plotting linearized response:

$$Z_i = \mathbf{X}_i' \hat{\beta} + (Y_i - \hat{\lambda}) / \hat{\lambda}$$

```
lambda<- predict(fit1, type="response")  
z <- predict(fit1) + (gala$Species-lambda)/lambda  
scatter.smooth(log(gala$Area), z,  
               ylab="Linearized response", main="(d)")
```

- $\log$  and the linearized response are similar

# Consider $\log(\text{Area})$



- Log transformation appears appropriate

# Residuals of Original and Log-Transformed Predictors

- Apply log transform to all predictors

```
>fit2 <- glm(Species~ log(Area)+log(Elevation)+  
  log(Nearest)+log(Scruz+0.1)+log(Adjacent),  
  family=poisson, data=gala)
```

```
>anova(fit1, fit2)
```

Analysis of Deviance Table



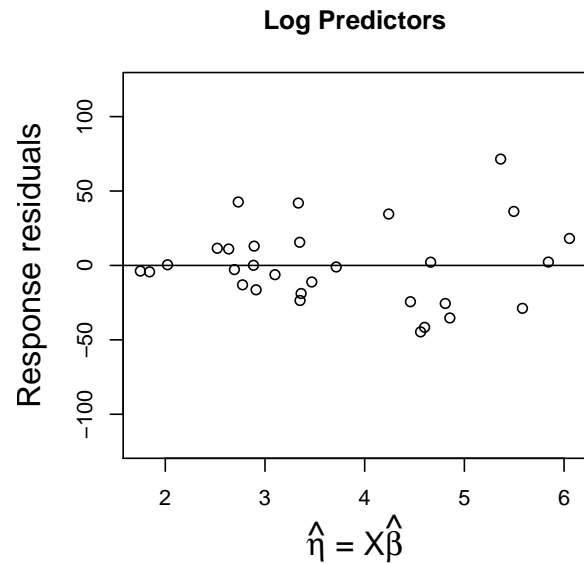

Model 1: Species ~ Area + Elevation +  
 Nearest + Scruz + Adjacent

Model 2: Species ~ log(Area) + log(Elevation) +  
 log(Nearest) + log(Scruz + 0.1) + log(Adjacent)

	Resid. Df	Resid. Dev	Df	Deviance
1	24	716.85		
2	24	359.12	0	357.72

- obtain a substantial reduction in deviance
- the model with all the transformed predictors yields a better overall fit
- Next: visualize the change in plots of response and deviance residuals

# Residuals of Original and Log-Transformed Predictors



- More compact residuals after the transform

# Model Diagnostics: Partial Residuals and Link

(a) Partial contribution of a predictor:

Plot  $Z_i - \mathbf{X}_i' \hat{\beta} + X_{ij} \hat{\beta}_j$  vs  $X_j$

```
lambda <- predict(fit2, type="response")
u <- (gala$Species-lambda)/lambda +
      coef(fit2)[2]*log(gala$Area)
scatter.smooth(log(gala$Area), u,
               ylab="Partial residual", main="Partial residuals")
```


(b) Choice of link function:

Plot  $Z_i$  vs  $\mathbf{X}_i' \hat{\beta}$

```
z <- predict(fit2) + (gala$Species-lambda)/lambda
scatter.smooth(predict(fit2), z,
               ylab="Linearized response",
               main="Diagnostics for link")
```

- No major problems in this dataset
  - Points roughly forming straight lines indicate no gross departures from the assumptions
  - Possible two outliers

# Model Diagnostics: Partial Residuals and Link



- No major problems in this dataset
  - Points roughly forming straight lines indicate no gross departures from the assumptions
  - Two potential outliers

# Model Diagnostics: Outliers

- Same concepts as in linear and logistic regression
- Automated plots in R (next slide)

```
for (i in 1:4)
  plot(fit2, which=i)
```

- Manual plots in R (following slide)

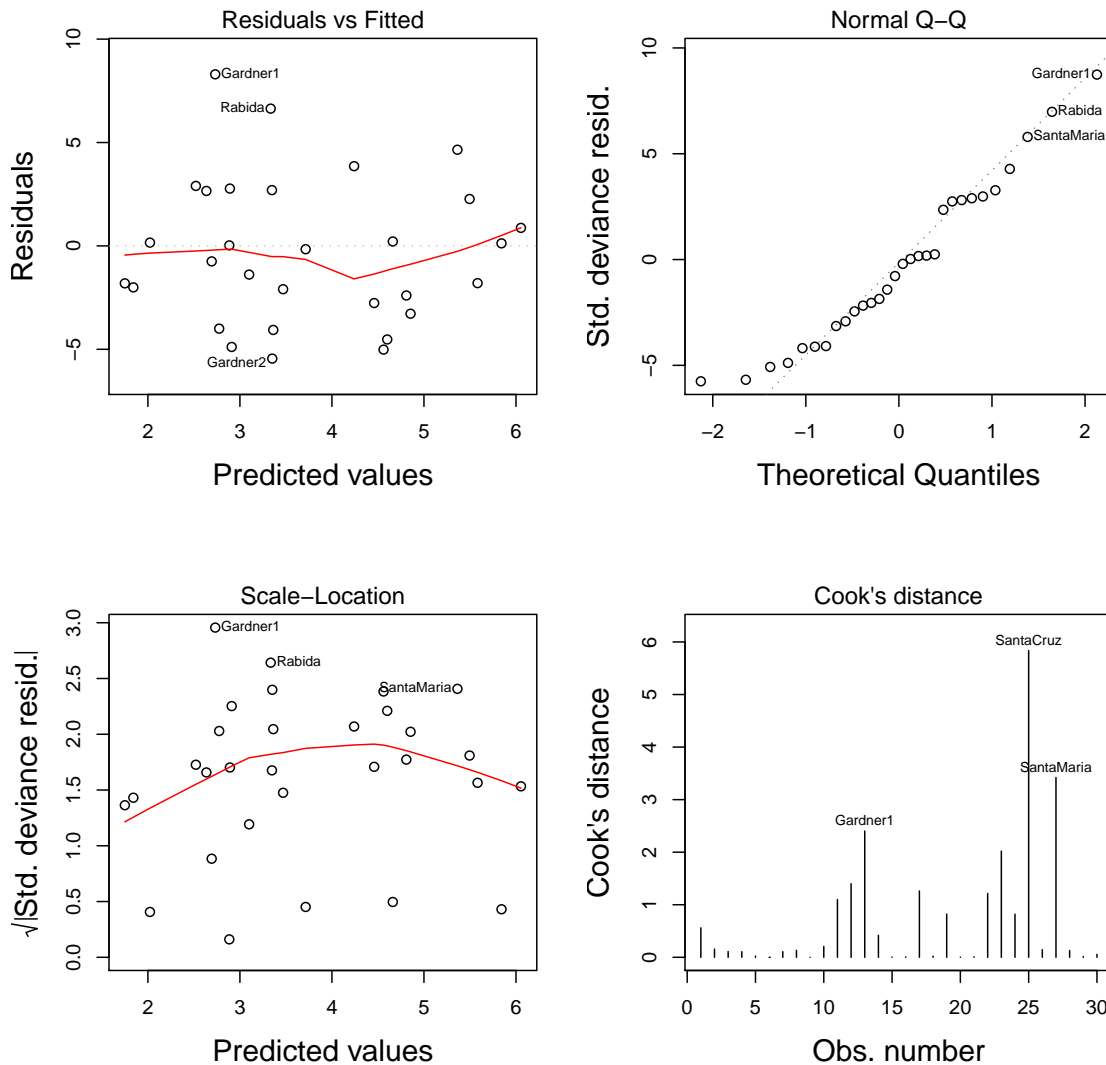
```
halfnorm( rstudent(fit2),
  main="Half-normal qqplot of Jackknife residuals")
```

```
halfnorm( influence(fit2)$hat,
  main="Half-normal qqplot of leverages")
```

```
halfnorm( cooks.distance(fit2),
  main="Half-normal qqplot of Cook's distance")
```

```
plot( influence(fit2)$coef[,5], xlab="Obs.No",
  ylab="dBeta",
  main="Change in Scruz coef after deletion")
```

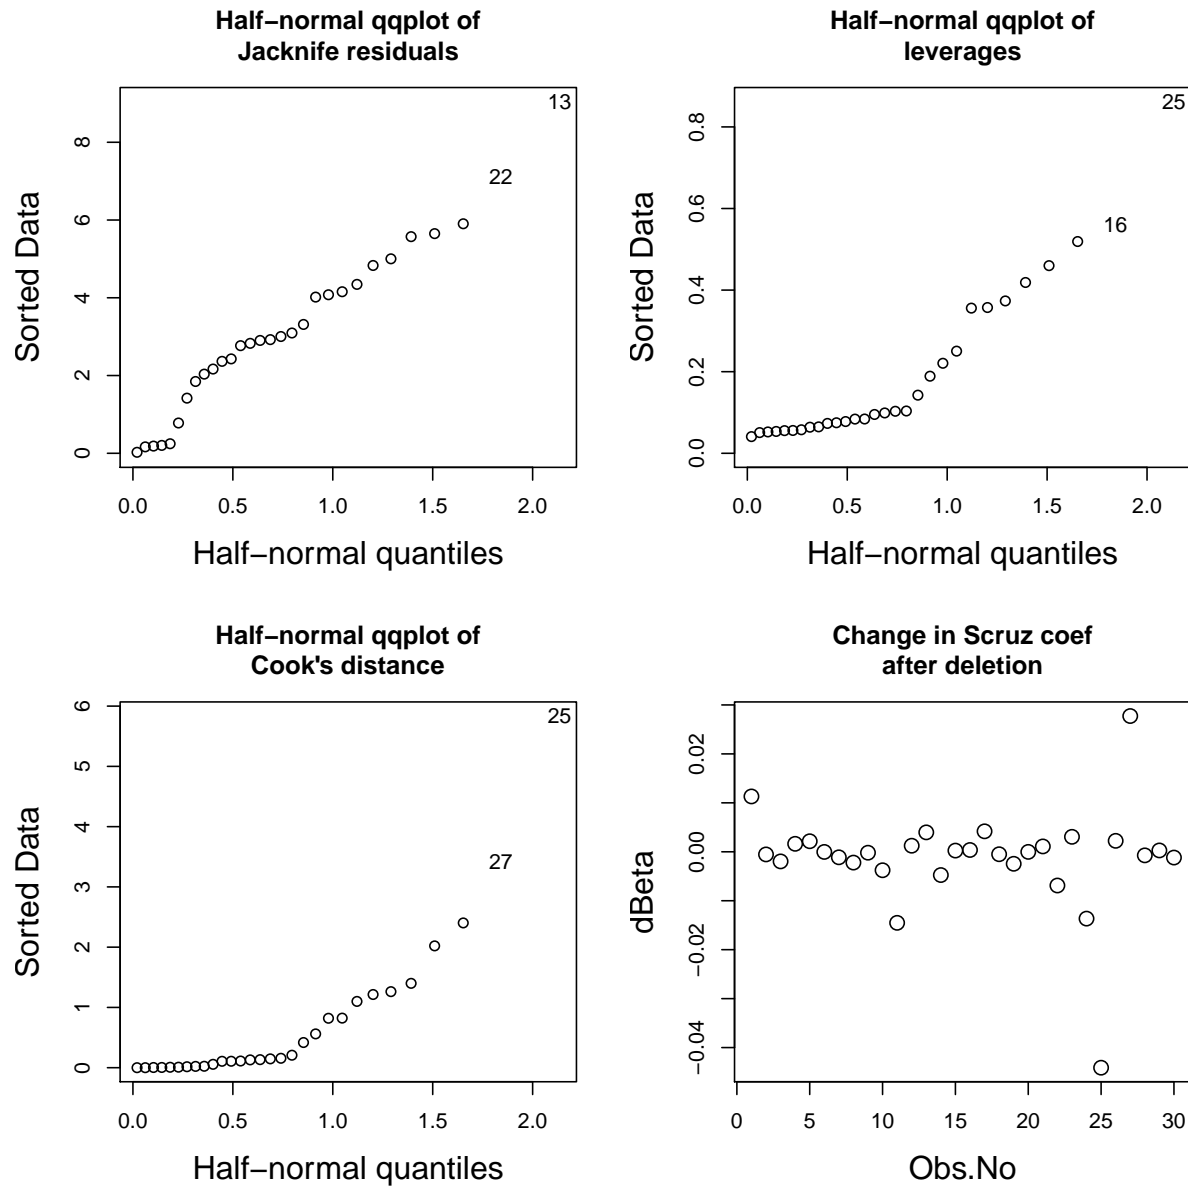
# Diagnostics: Outliers



- 'Residuals' = deviance residuals;
- 'Predicted values' =  $\mathbf{X}_i' \hat{\beta}$ ;
- 'Std. deviance resid' =  $r_{SD} = \frac{r_D}{\sqrt{\hat{\phi}(1-h_{ii})}}$  (p. 124);
- 'Cook's distance' (p. 125).



# Diagnostics: Outliers



- #25= 'Scrutz'. The influence can be due to adding 0.1 at the log transform.

# **Overdispersion (Faraway Ch. 3)**

# Overdispersion

- Suppose  $Y \sim \text{Poisson}(\lambda)$ 
  - as for Binomial response, overdispersion arises in case of hierarchical or clustered observations
- Suppose  $\lambda$  itself is a r.v. with  $E\{\lambda\} = \theta$ .
  - $E\{Y\} = E\{E\{Y|\lambda\}\} = \theta$
  - $Var\{Y\} = E\{Var\{Y|\lambda\}\} + Var\{E\{Y|\lambda\}\}$ 

$$= \theta + Var\{\lambda\} > \theta$$
- Quasipoisson distribution
  - $E\{Y\} = \lambda$
  - $Var\{Y\} = \phi\lambda$
  - Can estimate  $\hat{\phi}$ :  $\hat{\phi} = \frac{1}{n-p} X^2 = \frac{1}{n-p} \sum_{i=1}^n \frac{(y_i - \hat{\lambda}_i)^2}{\hat{\lambda}_i}$
- Comparing nested models in presence of overdispersion:
  - Test  $H_0$  : reduced model vs  $H_a$  : full model
  - $F = \frac{(D_0 - D_1)/(df_0 - df_1)}{\hat{\phi}_1} \overset{approx.}{\sim} F(df_0 - df_1, df_1)$

# Example: Galapagos Plants

- Fit a simplified model with 2 predictors

```
> fit3 <- glm(Species~ log(Area)+log(Adjacent),  
              family=poisson, data=gala)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	3.276683	0.044132	74.25	<2e-16 ***
log(Area)	0.375029	0.008023	46.74	<2e-16 ***
log(Adjacent)	-0.095747	0.006118	-15.65	<2e-16 ***

...


(Dispersion parameter for poisson family taken to be 1)

Null deviance: 3510.73 on 29 degrees of freedom  
Residual deviance: 395.54 on 27 degrees of freedom  
AIC: 562.37

- Log-log plot of mean vs variance (next slide)

```
lambda <- predict(fit3, type="response")  
plot(log(lambda),  
      log((gala$Species -lambda)^2),  
      xlab=expression(hat(lambda)),  
      ylab=expression((y - hat(lambda))^2))  
abline(0,1)
```

# Diagnostics: Overdispersion



- X axis:  $\log(E\{\widehat{Y}\})$  of each observation
- Y axis:  $\log(Var\{\widehat{Y}\})$  of each observation
- The variance is increasing with the mean
- The variance is larger than the mean  
→ indicates overdispersion

# Account for Overdispersion

- Estimate the dispersion parameter

```
> sum(residuals(fit3, type="pearson")^2) / fit3$df.res  
[1] 16.52745    # >> 1, i.e. overdispersion present
```

- Quasipoisson model

```
> summary( glm(Species~ log(Area)+log(Adjacent),  
              family=quasipoisson, data=gala) )
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	3.27668	0.17942	18.26	< 2e-16	***
log(Area)	0.37503	0.03262	11.50	6.56e-12	***
log(Adjacent)	-0.09575	0.02487	-3.85	0.000658	***
...					

(Dispersion parameter for quasipoisson family taken  
to be 16.52745)

Null deviance: 3510.73 on 29 degrees of freedom  
Residual deviance: 395.54 on 27 degrees of freedom  
AIC: NA

- Same parameters, larger SE, weaker significance

# **Rate Models (Faraway Ch. 3)**

# Motivation

- Each  $Y_i$  can represent a different interval in space or time
  - # of crimes in cities of different size
  - # of customers served by workers who work different hours
  - # cars running red light in different intersections
- Modeling  $Y_i$  as Poisson response is often more appropriate than Binomial
  - counts can be small as compared to the total (crimes)
  - the total may not be a count (workers, intersections)
- Goal: express a common effect of covariates on all  $Y_i$ , while accounting for differences in 'exposure'
  - 'exposure' needs to be a defined variable



# Formulation

- Model  $Y_i$  as

$$Y_i \stackrel{ind}{\sim} \text{Poisson}(\lambda_i), \quad \lambda_i = \text{exposure}_i e^{\mathbf{x}_i' \beta},$$

*i.e.*  $\log(\lambda_i) = \log(\text{exposure}_i) + \mathbf{x}_i' \beta$

- The model is equivalent to using *exposure* as a predictor with the coefficient=1
  - *exposure* is often called an *offset* in this context

## Example: Gamma Radiation (Faraway Sec. 3.2)

- Effect of gamma radiation on the numbers of chromosomal abnormalities (ca) in cells
- Different number of cells (cells, in hundreds) exposed to a dose and a rate of radiation
- cells can be viewed as 'exposure' variable


```
library(faraway)
data(dicentric)
?dicentric
```

```
> head(dicentric)
```

	cells	ca	doseamt	doserate
1	478	25	1	0.10
2	1907	102	1	0.25
3	2258	149	1	0.50
4	2329	160	1	1.00
5	1238	75	1	1.50
6	1491	100	1	2.00

# Visualize the data

- Interaction plot: proportion of cells with abnormality
    - The effect of dose rate is multiplicative
- ```
> with(dicentric,  
      interaction.plot(doseamt, doserate, ca/cells))
```



# Model Without Offset

```
> # indicate that 'dose' is categorical:
> dicentric$doseF <- factor(dicentric$doseamt)

> # the total number of cells as an independent predictor
> fit4 <- glm(ca ~ log(cells) + log(doserate)*doseF,
              family=poisson, data=dicentric)
```

...

Coefficients:

|                        | Estimate | Std. Error | z value | Pr(> z ) |     |
|------------------------|----------|------------|---------|----------|-----|
| (Intercept)            | -2.76534 | 0.38116    | -7.255  | 4.02e-13 | *** |
| log(cells)             | 1.00252  | 0.05137    | 19.517  | < 2e-16  | *** |
| log(doserate)          | 0.07200  | 0.03547    | 2.030   | 0.042403 | *   |
| doseF2.5               | 1.62984  | 0.10273    | 15.866  | < 2e-16  | *** |
| doseF5                 | 2.76673  | 0.12287    | 22.517  | < 2e-16  | *** |
| log(doserate):doseF2.5 | 0.16111  | 0.04837    | 3.331   | 0.000866 | *** |
| log(doserate):doseF5   | 0.19316  | 0.04299    | 4.493   | 7.03e-06 | *** |

...

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 916.127 on 26 degrees of freedom  
Residual deviance: 21.748 on 20 degrees of freedom  
AIC: 211.15

- The coefficient of log(cells) is close to 1, and the offset is appropriate

# Model With Offset

```
> fit5 <- glm(ca ~ offset(log(cells)) + log(doserate)*doseF,  
  family=poisson, data=dicentric)
```

...

Coefficients:

|                        | Estimate | Std. Error | z value | Pr(> z ) |     |
|------------------------|----------|------------|---------|----------|-----|
| (Intercept)            | -2.74671 | 0.03426    | -80.165 | < 2e-16  | *** |
| log(doserate)          | 0.07178  | 0.03518    | 2.041   | 0.041299 | *   |
| doseF2.5               | 1.62542  | 0.04946    | 32.863  | < 2e-16  | *** |
| doseF5                 | 2.76109  | 0.04349    | 63.491  | < 2e-16  | *** |
| log(doserate):doseF2.5 | 0.16122  | 0.04830    | 3.338   | 0.000844 | *** |
| log(doserate):doseF5   | 0.19350  | 0.04243    | 4.561   | 5.1e-06  | *** |

...

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 4753.00 on 26 degrees of freedom  
Residual deviance: 21.75 on 21 degrees of freedom  
AIC: 209.16

- The coefficients did not change much from the previous model.
- Good residual deviance, no evidence of overdispersion.

# Grouped vs Ungrouped Data

- Since  $Y_i \overset{ind}{\sim} \text{Poisson}(\lambda_i)$ ,  $i = 1, \dots$ ,  
 $\rightarrow \sum_i Y_i \sim \text{Poisson}(\sum_i \lambda_i)$ ,  
we can model aggregate data
  - as in Binomial or Multinomial response, log-likelihood only involves sums of  $y_i$  with same covariate patterns
  - e.g. add total cells and counts of abnormalities for entries with same doseamt and doserate
- Rate models: individual vs grouped
  - different deviances
  - same parameter estimates
  - same comparison of nested models
- Models with no offset: indiv. vs grouped
  - different deviances
  - same parameters except intercept
  - same comparison of nested models

**Allowing for Extra Variation:**

**Negative Binomial  
(Faraway Sec. 3.3)**

# Negative Binomial: Distribution

- Independent trials with  $P\{success\} = p$ 
  - $Z =$  the number of trials until  $k$ th success
  - $Z \sim NB(p, k), Z = k, k + 1, \dots$
- Probability distribution
  - $P\{Z = z\} = \binom{z-1}{k-1} p^k (1-p)^{z-k}$
  - $E\{Z\} = \frac{k}{p}$
  - $Var\{Z\} = \frac{k(1-p)}{p^2}$
- Obtained as the marginal distribution of  $Y|\lambda \sim Poisson(\lambda)$  where  $\lambda \sim G(k, \alpha)$ 
  - $P\{Y = y\} = \frac{1}{y!} \lambda^y e^{-\lambda}, Y = 0, 1, \dots$
  - $f(\lambda) = \frac{\alpha^k}{\Gamma(k)} \lambda^{k-1} e^{-\alpha\lambda}, \lambda > 0$



# Negative Binomial: Genesis

- Suppose  $Y|\lambda \sim \text{Poisson}(\lambda)$ ,  $\lambda \sim G(k, \alpha)$

– Joint distribution of  $Y$  and  $\lambda$ :

$$p(Y = y, \lambda) = \frac{\alpha^k}{\Gamma(k) y!} \lambda^{y+k-1} e^{-(\alpha+1)\lambda}$$

– Marginal distribution of  $Y$  (by integrating  $\lambda$ ):

$$\begin{aligned} p(Y = y) &= \frac{\alpha^k}{\Gamma(k) y!} \int_0^{\infty} \lambda^{y+k-1} e^{-(\alpha+1)\lambda} d\lambda \\ &= \frac{\alpha^k}{\Gamma(k) y!} \frac{\Gamma(y+k)}{(\alpha+1)^{y+k}} \\ &= \frac{\Gamma(y+k)}{\Gamma(k)\Gamma(y+1)} \left(\frac{\alpha}{\alpha+1}\right)^k \left(\frac{1}{\alpha+1}\right)^y \\ &= \frac{(y+k-1)!}{(k-1)!y!} \left(\frac{\alpha}{\alpha+1}\right)^k \left(\frac{1}{\alpha+1}\right)^y \\ &= \binom{y+k-1}{k-1} \left(\frac{\alpha}{\alpha+1}\right)^k \left(\frac{1}{\alpha+1}\right)^y \end{aligned}$$

- $\Rightarrow Z = Y + k \sim \text{NB}(\frac{\alpha}{\alpha+1}, k), Y = 0, 1, \dots$

# Negative Binomial: Model

- $Y + k \sim NB(\frac{\alpha}{\alpha+1}, k), Y = 0, 1, \dots$

$$- E\{Y\} = E\{Y + k\} - k = \frac{k}{\alpha/(\alpha+1)} - k = \frac{k}{\alpha} \stackrel{\text{denote}}{=} \mu$$

$$\begin{aligned} - Var\{Y\} &= Var\{Y + k\} = \frac{k(1-\alpha/(\alpha+1))}{\alpha^2/(\alpha+1)^2} = \frac{k}{\alpha} + \frac{k}{\alpha^2} \\ &= \mu + \mu^2/k \quad (= \text{Poisson variance} + \text{extra component}) \end{aligned}$$

- Log-likelihood of  $y_1, \dots, y_n$

$$\begin{aligned} & \sum_{i=1}^n \left( y_i \log \frac{1}{\alpha_i + 1} + k \log \frac{\alpha_i}{\alpha_i + 1} \right) + \text{constant} \\ & \stackrel{k/\alpha=\mu}{=} \sum_{i=1}^n \left( y_i \log \frac{\mu_i}{k + \mu_i} + k \log \frac{1}{k + \mu_i} \right) + \text{constant} \end{aligned}$$

$$- \text{Can model } \log \frac{\mu_i}{k + \mu_i} = \mathbf{x}_i' \beta$$

- $k$  is usually assumed same for all  $y_i$ 
  - fixed *a priori*, or estimated from the data
  - assumes constant coefficient of variation of  $\lambda$ :  
 $\sqrt{Var\{\lambda\}}/E\{\lambda\} = 1/\sqrt{k}$

# Example: Galapagos Plants

- Negative Binomial in an instance of the Exponential Family of distributions when  $k$  is fixed
- Fix  $k$ , use glm.

```
> fit.nb <- glm(Species~ log(Area)+log(Adjacent),  
  family=negative.binomial(1), data=gala)
```

Coefficients:

|               | Estimate | Std. Error | t value | Pr(> t ) |     |
|---------------|----------|------------|---------|----------|-----|
| (Intercept)   | 3.27257  | 0.15304    | 21.384  | < 2e-16  | *** |
| log(Area)     | 0.35100  | 0.03773    | 9.304   | 6.52e-10 | *** |
| log(Adjacent) | -0.03204 | 0.04015    | -0.798  | 0.432    |     |
| ...           |          |            |         |          |     |

(Dispersion parameter for Negative Binomial(1) family  
taken to be 0.4650222)

Null deviance: 54.069 on 29 degrees of freedom  
Residual deviance: 13.965 on 27 degrees of freedom  
AIC: 292.97

# Example: Galapagos Plants

- Use `glm.nb`, and estimate  $k$  from data

```
> fit.nb1 <- glm.nb(Species~ log(Area)+log(Adjacent),  
  data=gala)
```

Coefficients:

|               | Estimate | Std. Error | z value | Pr(> z )   |
|---------------|----------|------------|---------|------------|
| (Intercept)   | 3.27777  | 0.14495    | 22.613  | <2e-16 *** |
| log(Area)     | 0.34973  | 0.03541    | 9.875   | <2e-16 *** |
| log(Adjacent) | -0.03316 | 0.03737    | -0.887  | 0.375      |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(2.6196)  
family taken to be 1)

Null deviance: 134.240 on 29 degrees of freedom  
Residual deviance: 32.741 on 27 degrees of freedom  
AIC: 284.99

Theta: 2.620  
Std. Err.: 0.753

- Not a big difference in model fit in this example